

Data-Driven Device Failure Prediction Thesis Defense

Paul L. Jordan

Thesis Advisor
G.L. Peterson

Committee Members
M.J. Mendenhall, A.C. Lin, and
A.J. Sellers

18 August 2016



Overview

The AFIT of Today is the Air Force of Tomorrow.

Introduction

Related Work

Methodology

Experimental Results

Fault Injection

Under-Resourced Central Processing Unit (CPU)

Memory Corruption

Memory Leak

Conclusions & Future Work



Problem

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Computer systems fail, not often, but could have devastating consequences
- ▶ Redundancy can help, but it is expensive, complex, and only masks the root cause
- ▶ Survey paper of techniques for predicting failure using machine learning [1]
- ▶ Difficulty is in finding or obtaining training data



The AFIT of Today is the Air Force of Tomorrow.

- ▶ Adaptive Failure Prediction (AFP) framework [2]
- ▶ AFP wasn't capable of running on modern operating system
- ▶ AFP didn't exhaustively emulate all possible/realistic faults [3]



Solution

The AFIT of Today is the Air Force of Tomorrow.

- ▶ This work presents an extended AFP
- ▶ New realistic fault loads:
 - ▶ Memory Corruption
 - ▶ CPU Limitation
 - ▶ Memory Limitation (due to leak)
- ▶ Modernized fault injection tool: Windows Software Fault Injection Tool (W-SWFIT)
- ▶ Workload Generator: Distributed PowerShell Load Generator (D-PLG) [4]



- 7



The AFIT of Today is the Air Force of Tomorrow.




Figure: The timeline for OFP [1].

- ▶ Present Time: t
- ▶ Lead Time: Δt_l , is the total time at which a predictor makes an assessment about the current state.
- ▶ Data Window: Δt_d , represents the time window of data used for a predictor to make its assessment.
- ▶ Minimal Warning Time: Δt_w , is the amount of time required to avoid a failure if one is predicted.
- ▶ Prediction Period: Δt_p , is the time for which a prediction is valid.



Faults, Errors, and Failures

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Failure: Delivered service deviates from correct service
- ▶ Error: The point when things go wrong (Detected vs. Undetected)
- ▶ Fault: Hypothesized root cause of an error



Faults, Errors, and Failures

The AFIT of Today is the Air Force of Tomorrow.

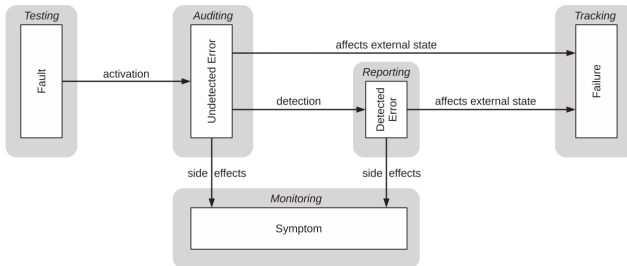


Figure: How faults and errors evolve into failure with the associated methods for detection represented by enclosing gray boxes [1].



Base Adaptive Failure Prediction (AFP)



The AFIT of Today is the Air Force of Tomorrow.

- ▶ Virtually clones target system
- ▶ Generates realistic workload for target to accomplish
- ▶ Synthesizes realistic faults that lead to failure
- ▶ Capture data from system to train failure prediction model



Base Adaptive Failure Prediction (AFP)

The AFIT of Today is the Air Force of Tomorrow.




Figure: The AFP framework [2].



Adaptive Failure Prediction (AFP)

Implementation

The AFIT of Today is the Air Force of Tomorrow.

Table: Hypervisor 1 configuration (sandbox/target).

Qty.	Role	Operating System	CPU / Mem.
1	DC	Win. Server 2008 R2	2 / 2 GB
1	Web	Win. Server 2008 R2	2 / 2 GB
5	Client	Win. 7	1 / 512 MB

Table: Hypervisor 2 configuration (controller).

Qty.	Role	Operating System	CPU / Mem.
1	RDP	Win. Server 2008 R2	1 / 4 GB
1	Log	Ubuntu 14.04 LTS	1 / 1 GB



Adaptive Failure Prediction (AFP) Implementation

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Workload Generator: D-PLG with five client machines
- ▶ Fault-Load Generator:
 - ▶ W-SWFIT
 - ▶ Memory Corruption
 - ▶ CPU Limitation
 - ▶ Memory Limitation (due to leak)
- ▶ Events Manager: rsyslog server with SolarWinds syslog forwarder
- ▶ Prediction Model: Support Vector Machine (SVM) and boosted decision trees in R



- ▶ Four fault loads tested on a Windows Domain Controller (DC) and an Apache web server:
 - ▶ W-SWFIT
 - ▶ Memory Corruption
 - ▶ CPU Limitation
 - ▶ Memory Limitation (due to leak)



Fault Injection

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Target process crashes immediately
- ▶ DC: restarts the computer
- ▶ Apache: starts a new child server process or parent process halts
- ▶ No indicators to use to train machine learning algorithm



Under-Resource CPU

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Third party application consumed all CPU time
- ▶ Virtual Machine (VM) resources were reduced
- ▶ Results
 - ▶ Both cases resulted in same behavior
 - ▶ Slower response times for both the DC and Apache web server
 - ▶ Target process would not fail



Memory Corruption

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Different from fault-injection in that it corrupts heap-space instead of program memory
- ▶ DC: corrupted the user database
- ▶ Apache: corrupted web content
- ▶ Results
 - ▶ Same as fault injection: either would not fail, or would crash immediately with no warning signs



Memory Leak

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Third party application consumed all available memory
- ▶ Only fault load that caused failure with indicators present in log messages prior to failure
- ▶ Trained two statistical models (SVM, and Boosted Decision Trees)
- ▶ As expected, both predictors performed adequately before software update, then poorly after
- ▶ After re-training with newly generated data performance once again was adequate




The AFIT of Today is the Air Force of Tomorrow.

- ▶ What is adequate?
 - ▶ Naïve predictor predicts non-failure prone at all times
 - ▶ Currently no form of prediction is taking place in operational environment
 - ▶ Machine learning classification algorithms evaluated using confusion matrix at best F-Measure, Receiver Operating Characteristic (ROC), Area Under the Curve (AUC), and Precision/Recall Curves [1, 5]




Support Vector Machine (SVM) Performance

The AFIT of Today is the Air Force of Tomorrow.



(a) Precision/Recall Curve.



(b) ROC Curve (AUC = 0.8664).

Figure: Test data performance of the SVM prediction method on failure data obtained by consuming all available memory until target application fails.

- ▶ Before software update, AUC: 0.8664
- ▶ After re-training, AUC: N/A (highest F-Measure: 0.4380)




Table: Confusion matrix on test data created before software updates on threshold with highest F-Measure (0.8739) using SVM.

		Actual	
		Fail	No-Fail
Predicted	Fail	52	6
	No-Fail	9	607




Boosted Decision Tree Performance

The AFIT of Today is the Air Force of Tomorrow.



(a) Precision/Recall Curves.



(b) ROC Curves.

Figure: Performance of the boosting prediction method on data generated by consuming all available memory until target application fails.

- ▶ Before software update, AUC: 0.9984
- ▶ After software update, AUC: 0.4854
- ▶ After re-training, AUC: 0.9801



Boosted Decision Tree Performance I

The AFIT of Today is the Air Force of Tomorrow.

Table: Confusion matrix on test data created before software updates on threshold with highest F-Measure (0.9917) using boosting.

		Actual	
		Fail	No-Fail
Predicted	Fail	60	0
	No-Fail	1	412



Boosted Decision Tree Performance II

The AFIT of Today is the Air Force of Tomorrow.

Table: Post-update failure data confusion matrix on threshold with highest F-Measure (0.4691) using model trained on failure data generated before software update.

		Actual	
		Fail	No-Fail
Predicted	Fail	19	1
	No-Fail	42	222



		Actual	
		Fail	No-Fail
Predicted	Fail	58	5
	No-Fail	3	218



Conclusions

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Extended AFP presented can predict realistic failure in production systems
- ▶ Capable of adapting to underlying system changes
- ▶ Vulnerability to certain types of failure can come and go, consequently, all fault loads should be used



Future Work

The AFIT of Today is the Air Force of Tomorrow.

- ▶ Further explore how best to implement fault injection and its true impact
- ▶ Integrate real failure data
- ▶ Further validation and automation
- ▶ Implement and make operational



Summary

The AFIT of Today is the Air Force of Tomorrow.

Introduction

Related Work

Methodology

Experimental Results

Fault Injection

Under-Resourced CPU

Memory Corruption

Memory Leak

Conclusions & Future Work



The AFIT of Today is the Air Force of Tomorrow.

Questions?



References I

The AFIT of Today is the Air Force of Tomorrow.

- [1] F. Salfner, M. Lenk, and M. Malek, "A survey of online failure prediction methods," *ACM Computing Surveys (CSUR)*, vol. 42, no. 3, 2010.
- [2] I. Irrera, M. Vieira, and J. Duraes, "Adaptive failure prediction for computer systems: A framework and a case study," in *Proceedings of the 2015 IEEE 16th International Symposium on High Assurance Systems Engineering (HASE 2015)*, pp. 142–149, 2015.
- [3] N. Kikuchi, T. Yoshimura, R. Sakuma, and K. Kono, "Do injected faults cause real failures? a case study of linux," in *Proceedings of the 25th IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW 2014)*, pp. 174–179, 2014.



References II

The AFIT of Today is the Air Force of Tomorrow.

- [4] P. Jordan, C. Van Patten, G. Peterson, and A. Sellers, “Distributed powershell load generator (D-PLG): A new tool for dynamically generating network traffic,” in *Proceedings of the 6th International Conference on Simulation and Modeling Methodologies, Technologies, and Applications (SIMULTECH 2016)*, pp. 195–202, July 2016.
- [5] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An Introduction to Statistical Learning: With Applications in R*. Springer Publishing Company, Incorporated, 2014.