

Air Force Institute of Technology



The AFIT of Today is the Air Force of Tomorrow.



Data-Driven Device Failure Prediction

Paul I. Jordan

Thesis Advisor G.L. Peterson

Committee Members
M.J. Mendenhall, A.C. Lin, and
A.J. Sellers
18 August 2016



Introduction

Related Work

Methodology

Experimental Results

Fault Injection Under-Resourced Central Processing Unit (CPU) Memory Corruption Memory Leak

Conclusions & Future Work







- Computer systems fail, not often, but could have devastating consequences
- Redundancy can help, but it is expensive, complex, and only masks the root cause
- Survey paper of techniques for predicting failure using machine learning [1]
- Difficulty is in finding or obtaining training data





- ▶ Adaptive Failure Prediction (AFP) framework [2]
- ► AFP wasn't capable of running on modern operating system
- ► AFP didn't exhaustively emulate all possible/realistic faults [3]





- ► This work presents an extended AFP
- New realistic fault loads:
 - Memory Corruption
 - CPU Limitation
 - Memory Limitation (due to leak)
- Modernized fault injection tool: Windows Software Fault Injection Tool (W-SWFIT)
- Workload Generator: Distributed PowerShell Load Generator (D-PLG) [4]





Related Work



Proactive Fault Management (PFM)



The AFIT of Today is the Air Force of Tomorrow.

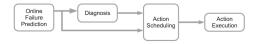


Figure: The stages of proactive fault management [1].

- Online Failure Prediction (OFP)
- Diagnosis
- Action Scheduling
- Action Execution



Online Failure Prediction (OFP)



The AFIT of Today is the Air Force of Tomorrow.

 Act of evaluating a running system in real time to make a prediction about whether a failure in a future state is imminent [1]



Online Failure Prediction (OFP)



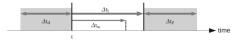


Figure: The timeline for OFP [1].

- Present Time: t
- ▶ Data Window: Δt_d , represents the time window of data used for a predictor to make its assessment.
- ▶ Lead Time: Δt_I , represents the time between when failure is predicted and when that failure will occur.
- Minimal Warning Time: Δt_w , is the amount of time required to avoid a failure if one is predicted.
- ▶ Prediction Period: Δt_p , is the time for which a prediction is valid.



Faults, Errors, and Failures



- ▶ Failure: Delivered service deviates from correct service
- Error: The point when things go wrong (Detected vs. Undetected)
- Fault: Hypothesized root cause of an error



Faults, Errors, and Failures



The AFIT of Today is the Air Force of Tomorrow.

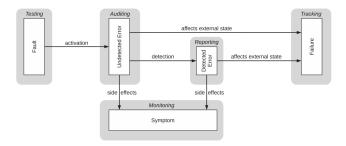


Figure: How faults and errors evolve into failure with the associated methods for detection represented by enclosing gray boxes [1].



Base Adaptive Failure Prediction (AFP)



- Virtually clones target system
- Generates realistic workload for target to accomplish
- Synthesizes realistic faults that lead to failure
- Capture data from system to train failure prediction model



Base Adaptive Failure Prediction (AFP)



The AFIT of Today is the Air Force of Tomorrow.




Figure: The AFP framework [2].





Methodology



Adaptive Failure Prediction (AFP) Implementation



The AFIT of Today is the Air Force of Tomorrow.

Table: Hypervisor 1 configuration (sandbox/target).

Qty.	Role	Operating System	CPU / Mem.
1	DC	Win. Server 2008 R2	2 / 2 GB
1	Web	Win. Server 2008 R2	2 / 2 GB
5	Client	Win. 7	1 / 512 MB

Table: Hypervisor 2 configuration (controller).

Qty.	Role	Operating System	CPU / Mem.
1	RDP	Win. Server 2008 R2	1 / 4 GB
1	Log	Ubuntu 14.04 LTS	1 / 1 GB



Adaptive Failure Prediction (AFP) Implementation



- Workload Generator: D-PLG with five client machines
- Fault-Load Generator:
 - ▶ W-SWFIT
 - Memory Corruption
 - CPU Limitation
 - Memory Limitation (due to leak)
- Events Manager: rsyslog server with SolarWinds syslog forwarder
- Prediction Model: Support Vector Machine (SVM) and boosted decision trees in R





Experimental Results







- Four fault loads tested on a Windows Domain Controller (DC) and an Apache web server:
 - W-SWFIT
 - Memory Corruption
 - CPU Limitation
 - Memory Limitation (due to leak)



Fault Injection



- ► Target process crashes immediately
- ▶ DC: restarts the computer
- ► Apache: starts a new child server process or parent process halts
- ▶ No indicators to use to train machine learning algorithm



Under-Resource CPU



- Third party application consumed all CPU time
- Virtual Machine (VM) resources were reduced
- Results
 - Both cases resulted in same behavior
 - ▶ Slower response times for both the DC and Apache web server
 - Target process would not fail



Memory Corruption



- ▶ Different from fault-injection in that it corrupts heap-space instead of program memory
- DC: corrupted the user database
- Apache: corrupted web content
- Results
 - Same as fault injection: either would not fail, or would crash immediately with no warning signs



Memory Leak



- Third party application consumed all available memory
- Only fault load that caused failure with indicators present in log messages prior to failure
- Trained two statistical models (SVM, and Boosted Decision Trees)
- As expected, both predictors performed adequately before software update, then poorly after
- ► After re-training with newly generated data performance once again was adequate



Memory Leak



The AFIT of Today is the Air Force of Tomorrow.

What is adequate?

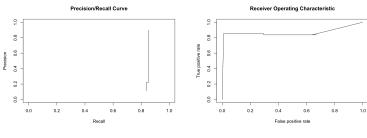
- ▶ Naïve predictor predicts non-failure prone at all times
- Currently no form of prediction is taking place in operational environment
- ► Machine learning classification algorithms evaluated using confusion matrix at best F-Measure, Receiver Operating Characteristic (ROC), Area Under the Curve (AUC), and Precision/Recall Curves [1, 5]



Support Vector Machine (SVM) Performance



The AFIT of Today is the Air Force of Tomorrow.



- (a) Precision/Recall Curve.
- (b) ROC Curve (AUC = 0.8664).

Figure: Test data performance of the SVM prediction method on failure data obtained by consuming all available memory until target application fails.

- ▶ Before software update, AUC: 0.8664
- ► After re-training, AUC: N/A (highest F-Measure: 0.4380)



Support Vector Machine (SVM) Performance



The AFIT of Today is the Air Force of Tomorrow.

Table: Confusion matrix on test data created before software updates on threshold with highest F-Measure (0.8739) using SVM.

Actual

	Fail	No-Fail
Fail	52	6
No-Fail	9	607



Boosted Decision Tree Performance



The AFIT of Today is the Air Force of Tomorrow.

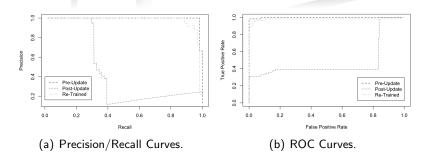


Figure: Performance of the boosting prediction method on data generated by consuming all available memory until target application fails.

- Before software update, AUC: 0.9984
- After software update, AUC: 0.4854
- After re-training, AUC: 0.9801



Boosted Decision Tree Performance I



The AFIT of Today is the Air Force of Tomorrow.

Table: Confusion matrix on test data created before software updates on threshold with highest F-Measure (0.9917) using boosting.

Actual

	Fail	No-Fail
Fail	60	0
No-Fail	1	412



Boosted Decision Tree Performance II



The AFIT of Today is the Air Force of Tomorrow.

Table: Post-update failure data confusion matrix on threshold with highest F-Measure (0.4691) using model trained on failure data generated before software update.

Actual

	Fail	No-Fail
Fail	19	1
No-Fail	42	222



Boosted Decision Tree Performance III



The AFIT of Today is the Air Force of Tomorrow.

Table: Post-update failure data confusion matrix on threshold with highest F-Measure (0.9355) using model trained on failure data generated after software update.

Actual

	Fail	No-Fail
Fail	58	5
No-Fail	3	218





Conclusions & Future Work





- Extended AFP presented can predict realistic failure in production systems
- Capable of adapting to underlying system changes
- Vulnerability to certain types of failure can come and go, consequently, all fault loads should be used





- ► Further explore how best to implement fault injection and its true impact
- Integrate real failure data
- Further validation and automation
- Implement and make operational



Summary



The AFIT of Today is the Air Force of Tomorrow.

Introduction

Related Work

Methodology

Experimental Results

Fault Injection Under-Resourced CPU Memory Corruption Memory Leak

Conclusions & Future Work





Questions?







- [1] F. Salfner, M. Lenk, and M. Malek, "A survey of online failure prediction methods," *ACM Computing Surveys (CSUR)*, vol. 42, no. 3, 2010.
- [2] I. Irrera, M. Vieira, and J. Duraes, "Adaptive failure prediction for computer systems: A framework and a case study," in Proceedings of the 2015 IEEE 16th International Symposium on High Assurance Systems Engineering (HASE 2015), pp. 142–149, 2015.
- [3] N. Kikuchi, T. Yoshimura, R. Sakuma, and K. Kono, "Do injected faults cause real failures? a case study of linux," in Proceedings of the 25th IEEE International Symposium on Software Reliability Engineering Workshops (ISSREW 2014), pp. 174–179, 2014.







- [4] P. Jordan, C. Van Patten, G. Peterson, and A. Sellers, "Distributed powershell load generator (D-PLG): A new tool for dynamically generating network traffic," in *Proceedings of* the 6th International Conference on Simulation and Modeling Methodologies, Technologies, and Applications (SIMULTECH 2016), pp. 195–202, July 2016.
- [5] G. James, D. Witten, T. Hastie, and R. Tibshirani, An Introduction to Statistical Learning: With Applications in R. Springer Publishing Company, Incorporated, 2014.