



Data Driven Device Failure Prediction

THESIS

Paul L. Jordan, 1st Lt, USAF
AFIT/GCS/ENG/17-M

**DEPARTMENT OF THE AIR FORCE
AIR UNIVERSITY**

AIR FORCE INSTITUTE OF TECHNOLOGY

Wright-Patterson Air Force Base, Ohio

DISTRIBUTION STATEMENT A
APPROVED FOR PUBLIC RELEASE; DISTRIBUTION UNLIMITED.

The views expressed in this document are those of the author and do not reflect the official policy or position of the United States Air Force, the United States Department of Defense or the United States Government. This material is declared a work of the U.S. Government and is not subject to copyright protection in the United States.

AFIT/GCS/ENG/17-M

DATA DRIVEN DEVICE FAILURE PREDICTION

THESIS

Presented to the Faculty
Department of Electrical and Computer Engineering
Graduate School of Engineering and Management
Air Force Institute of Technology
Air University
Air Education and Training Command
in Partial Fulfillment of the Requirements for the
Degree of Master of Science in Computer Science

Paul L. Jordan, B.S.C.S.

1st Lt, USAF

April 15, 2016

DISTRIBUTION STATEMENT A
APPROVED FOR PUBLIC RELEASE; DISTRIBUTION UNLIMITED.

AFIT/GCS/ENG/17-M

DATA DRIVEN DEVICE FAILURE PREDICTION

THESIS

Paul L. Jordan, B.S.C.S.
1st Lt, USAF

Committee Membership:

Dr. M. J. Mendenhall
Chair

Dr. G. L. Peterson
Member

Maj A. C. Lin, PhD
Member

Maj A. J. Sellers, PhD
Member

Abstract

This is the start of my abstract that I will write later.

Acknowledgements

I would like to thank...

Paul L. Jordan

Table of Contents

	Page
Abstract	iv
Acknowledgements	v
List of Figures	viii
List of Tables	ix
I. Introduction	1
1.1 Problem Statement	2
1.2 Impact of Research	4
1.3 Assumptions and Limitations	4
1.4 Expected Results of Research	4
II. Overview of Online Failure Prediction	6
2.1 Background	6
Definitions	6
Measures of Performance	10
2.2 Approaches to Online Failure Prediction (OFP)	15
OFP Taxonomy	15
Data-Driven Online Failure Prediction	16
Industry Approaches to Online Failure Prediction	20
Adaptive Failure Prediction (AFP) Framework	20
2.3 Summary	22
III. Methodology	23
3.1 Failure Data Generation	23
Preparation Phase	24
Execution Phase	24
Training Phase	26
3.2 Implementation of the AFP	29
AFP Framework Implementation	29
AFP Modules	29
3.2.1 Controller Hypervisor	30
3.2.2 Sandbox Hypervisor	39
3.2.3 Target Hypervisor	41
3.3 Extensions to the AFP	41
IV. Experimental Results and Analysis	43

	Page
V. Conclusion and Future Work	44
Bibliography	45

List of Figures

Figure	Page
1	Proactive Fault Management [19]6
2	Failure Flow Diagram [19]9
3	Online Failure Prediction [19]9
4	Sample Precision/Recall Curves [19].....14
5	Sample ROC Plots [19]14
6	Pattern recognition in reported errors [19].....16
7	AFP Framework Implementation [13].....21
8	AFP Execution Phase [13]25
9	AFP Training Phase [13]27
10	Annotated AFP Framework [13]30
11	Domain Controller Packets per Second37
12	Client Packets per Second37
13	Test 1: Domain Controller Metrics38

List of Tables

Table		Page
1	Hypervisor 1 Configuration (Sandbox/Target)	31
2	Hypervisor 2 Configuration (Controller)	31
3	Table of Faults Injected [7]	34
4	Funtion Entry/Exit Patterns (IA32) [7]	34
5	Funtion Entry/Exit Patterns (x86-64) [7]	34

DATA DRIVEN DEVICE FAILURE PREDICTION

I. Introduction

We as humans have always shared a curiosity about the future. As we become more and more dependent upon computers as they permeate their way into our daily lives, the software that controls these machines is ultimately still subject to human error. Some of these system are responsible for sustaining human life and ultimately are not protected from that same human error. Consequently, under the wrong circumstances a machine failure could result in loss of human life. Critical infrastructure and Air Force missions depend on the reliability of computer systems. As a result, being able to predict pending failue in computer systems can offer tremendous and potentially life-saving applications in todays technologically advanced world. While actually being able to accurately predict failure has unfortunately not been proven possible, there has been an enormous amount of time and energy spent over the past several decades attempting to make educated predictions about the failure of machines through the use of sophisticated machine learning algorithms. In this research we explore the application of a new framework developed to automatically re-train machine learning based failure predictors. There a number of ways to reduce the number of errors produced by a piece of software, but the software development life-cycle is shrinking and less time and effort are being devoted to reducing errors before deployment. This leaves real-time error prevention or handling.

In recent years, it seems many of the cloud based computing companies have attempted to solve this problem by making all of their services massively redundant. As hardware becomes more affordable, this is an effective approach in many ways

but ultimately is still not cost efficient. In some cases, funds may not be available to achieve this sort of redundancy. Consequently, this research focuses on a small piece of the general field of reliable computing: online failure prediction (OFP). OFP is the act of attempting to predict when failures are likely so that they can be avoided. A great deal of work has been done in this field which we outline in chapter II, but much of it has gone unimplemented due to the complex and manual task of training a prediction model. If the underlying system changes at all (which in today's world is a common occurrence due to the aforementioned shrinking software development life-cycle) the efficacy of a prediction model can be drastically reduced if not rendered completely useless until it is retrained. This research explores an implementation of a new framework for automatically retraining a predictor after such an underlying system change. More specifically, we present our results after implementing this framework using a Microsoft Windows Server domain controller. We then apply successive software updates until the model we have selected becomes useless and allow the framework to re-train our predictor.

1.1 Problem Statement

Predicting and alerting on impending network service failures currently uses thresholds and rules on discrete items in enterprise system logs. For example, if the central processing unit (CPU) and memory usage on a device exceeds 90%, then an alert may be issued. This approach works, but only for certain types of failures and in order to minimize the false positives, it only makes recommendations minutes before a failure, or when the system is in an already degraded performance mode. To maintain network resilience, the operational organizations responsible for communications support desperately need some means of gaining lead-time before a service failure occurs.

Preceding a service failure event, multiple indicators spread disparate sources, perhaps over a long period of time, may appear in system logs. The log entries of interest are also quite rare compared with normal operations. Because of these constraints, identifying failure indicators can be nearly impossible for humans to perform. Further, in most cases, restoring service is more important than identifying the indicators that may or may not have existed.

Failure prediction can be approached in many ways. Arguably the simplest approach is to use everyday statistical analysis to, for example, determine the mean time between failures of specific components. The analysis of all components making up a system can be aggregated to make predictions about that system using a set of statistics-based or business-relevant rules. Unfortunately, the complexity of modern architectures has outpaced such off-line statistical-based analysis, which has driven the advancement of OFP. OFP differs from other means of failure prediction in that it focuses on classifying the current running state of a machine as either failure prone or not, or in such a way that it describes the confidence in how failure prone a system is at present [19].

Unfortunately, in recent years much of the work in OFP has gone unused due to the dramatic decrease in cost and complexity involved in building hardware-based redundant systems. Furthermore, in most cases OFP implements machine learning algorithms that require manual re-training after underlying system changes. More troubling is that these system changes are becoming more frequent as the software development life cycle moves toward a more continuous integration model. To help solve these challenges, the framework presented in [13] uses simulated faults to automatically re-train a prediction algorithm to make implementing OFP approaches easier. We propose to expand the work in [13] to capture developments since its writing and generalize it so it works for a broader class of devices.

1.2 Impact of Research

Every day, many of the Air Force's critical missions depend on our computer infrastructure. An essential piece of this infrastructure is the authentication mechanisms that protect our sensitive information. Unfortunately, the software at the core of this infrastructure is written and maintained by humans and thus susceptible human error. This research will enable the Air Force and many others that use the Microsoft Enterprise Infrastructure to accurately predict pending service outages thereby providing lead-time in order to avoid those outages. The result is cost savings in personnel, equipment, but isn't limited to cost savings. It is difficult to quantify the risk of mission failure due to network service outage.

1.3 Assumptions and Limitations

As previously stated, it has not been proven possible to accurately predict future events without a priori knowledge. As a result, this research depends on the fact that there are indicators of failure present and available to us with enough lead-time to accurately make decisions and take mitigation action should failure be predicted based on these indicators. This research presents a method of predicting failure, but this method is completely useless at predicting *act of God* events. Further, this method is capable of predicting failure based on software failures and cannot provide any useful information about malicious attacks against the target software.

1.4 Expected Results of Research

Because a prediction method is not presently deployed on any Air Force network, any level of dependable prediction will be better than what is currently available. This research will attempt to show that after an underlying system change, this method of

predicting failure will be capable of automatically training a more effective prediction algorithm so that this technique can be implemented on an Air Force network with little to no impact on manpower. Consequently, we expect that this research will inform decision makers and allow them to implement this technique in a production environment.

Specifically, we believe the technique presented in this research could most effectively be implemented and used by the Cyber Security and Control System (CSCS) weapon system employed at the 561st and 83d Network Operation Squadrons (NOS) and their associated detachments to reduce the number of network service outages, increasing uptime, leading to improved mission effectiveness in both the support and operational domains. Further, we believe this technique general enough to be employed outside of the Air Force to increase mission effectiveness across the Department of Defense (DOD). External to the DOD, this research further generalizes an approach that could be used to help increase availability of nearly any computer system.

II. Overview of Online Failure Prediction

This chapter reviews current research regarding online failure prediction and its many approaches to build a foundation for this research. Further, a taxonomy of approaches was developed here [19], this chapter updates that taxonomy and classifies approaches since its publication using it.

2.1 Background

In 2010, Salfner et al. published a survey paper that provides a comprehensive summary of the state of the art on the topic of OFP [19]. In addition to the review of the literature up to the point of publication, they provide a summary of definitions (see Section 2.1) and measures of performance (see Section 2.1) commonly used in the community for couching the OFP discussion.

Definitions.

Proactive Fault Management:.

Figure 1 shows the components in the PFM process, and in that process the final three stages deine how much lead time is required to avoid a failure when predicted during OFP. Lead time, although defined later, is a critical element in any approach to OFP.




Figure 1. The four stages of proactive fault management and how they integrate [19].

Although there is a lot of utility in studying PFM in its entirety, the crux to its success is at the beginning of this four stage process. This research focusses mainly on that first step however, we include here a brief overview of the remaining steps as they define how much *lead time* is required to avoid a predicted failure in order for the first step (OFP) to be of any use. *Lead time* is formally defined later in this section, but it should be noted now that it is a critical element of a failure prediction approach.

OFP is defined as the first step in PFM seen in Figure 1. The last three stages of PFM are diagnosis, action scheduling, and action execution. Once failure has been predicted, a fault tolerant system must determine what will cause the failure. This stage is called the *diagnosis* stage or “root-cause analysis” stage. During the *diagnosis* stage, the analysis must be conducted so that a system knows which remediation actions are possible. After it is determined what will cause a failure, a fault tolerant system must schedule a remediation action that is either performed by an operator or done automatically. This stage is known as the *action scheduling* stage and normally takes as input the cost of performing an action, confidence in prediction, effectiveness/complexity of remedy action and makes a decision about what action to perform based on that input. In some cases a remedy action can be so simple that even if the confidence in the prediction is low, the action can still be performed with little impact on the overall system and its users. A thorough analysis of the trade-off between cost of avoidance and confidence in prediction and the associated benefits is described in [3]. Finally, in order to avoid failure, a system must execute the scheduled remediation action or let an operator know which actions can be taken in a stage called *action execution*.

Faults, Errors, Symptoms, and Failures:.

As [19] points out, many attempts have been made to define faults, errors, symptoms, and failures. In this research we use the definitions from [1] as interpreted and extended in [19] for the following terms: failure; error (detected versus undetected); fault; and symptom.

Failure is an event that occurs when the delivered service deviates from correct service. In other words, things can go wrong internally; as long as the output of a system is what is expected, failure has not occurred. An *error* is the part of the total state of the system that may lead to its subsequent service failure. *Errors* are characterized as the point when things go wrong [19]. Fault tolerant systems can handle errors without necessarily evolving into failure. There are two kinds of errors. First, a *detected error* is an error that is reported to a logging service. In other words, if it can be seen in a log then it is a detected error. Second, *undetected errors* are errors that have not been identified by an error detector. Undetected errors are things like memory leaks. The error certainly exists, but as long as there is usable memory, it is not likely to be reported to a logging service. Once the system runs out of usable memory, undetected errors will likely show up in logs and become a detected errors. A *fault* is the hypothesized root cause of an error. Faults can remain dormant for some time before manifesting themselves and causing an incorrect system state. In the memory leak example, the missing *free* statement in the source code would be the fault. A *symptom* is an out-of-norm behavior of a system's parameters caused by errors, whether detected or undetected. In the memory leak example, a possible symptom of the error might be delayed response times due to sluggish performance of the overall system.




Figure 2. How faults and errors evolve into failure with the associated methods for detection represented by enclosing gray boxes [19].

Figure 2 illustrates the differences between faults, errors, symptoms and failures. It further shows how faults can evolve into failures and how specifically at each of the steps, they can be detected. We define these specifically because in Section 2.2 we outline a taxonomy of OFP approaches introduced in [19] which divides the OFP approaches into four major categories based on the four methods for detecting errors in running systems: auditing, reporting, monitoring, and tracking.




Figure 3. The timeline for OFP [19].

Figure 3 demonstrates the timeline associated with OFP. It does so based on a set of parameters drawn from the literature from the work in [19]. The parameters used by the community to define a predictor are as follows:

- Present Time: t

- Lead Time: Δt_l , is the total time at which a predictor makes an assessment about the current state.
- Data Window: Δt_d , represents the time from which data is used for a predictor uses to make its assessment.
- Minimal Warning Time: Δt_w , is the amount of time required to avoid a failure if one is predicted.
- Prediction Period: Δt_p , is the time for which a prediction is valid. As $\Delta t_p \rightarrow \infty$, the accuracy of the predictor approaches 100% because every system will eventually fail. As this happens, the usefulness of a predictor is diminished.

As the above parameters are adjusted, predictors can become more or less useful. For example, it is clear that as we look further into the future potentially increasing *lead time*, our confidence in our prediction is likely to be reduced. On the other hand, if *lead time* is too small, we may not be able to effectively implement a remediation action. In general, OFP approaches seek to find a balance between the parameters, within an acceptable bound depending on application, to achieve the best possible performance.

Measures of Performance.

In order to accurately compare OFP approaches, standard measures of performance are used. A widely accepted set of measures of performance used by the community is presented in [19]. Before we begin the outline of the measures of performance used to evaluate and compare failure prediction approaches, it is important to note that OFP is done based on statistical analysis of known data. In other words, in order to calculate the following outlined measures of performance, predictors must

be evaluated against labeled data sets. Typically, the labeled data set is divided into three parts:

1. Training Set: A data set that allows a prediction model to establish and optimize its parameters
2. Validation Set: The parameters selected in the training phase are then validated against a separate data set
3. Test Set: The predictor is finally run against a final previously unevaluated data set to assess generalizability

During the test phase, true positives (negatives) versus false positives (negatives) are determined in order to compute the measures of performance in this section. The following terms and associated abbreviations are used: *True Positive* (TP) is when failure has been predicted and then actually occurs; *False Positive* (FP) is when failure has been predicted and then does not occur; *True Negative* (TN) is when a state has been accurately classified as non-failure prone; *False Negative* is when a state has been classified as non-failure prone and a failure occurs.

Precision and Recall:.

Precision and recall are the most popular measures of performance used when presenting and comparing OFP approaches. The two are related and often times improving precision results in reduced recall. For example, if a predictor is more accurate with a prediction then that predictor may be less likely to classify a state as failure-prone in an effort to minimize false-positives.

Precision is the number of correctly identified failures over number of all predicted failures. In other words, it reports, out of the predictions of a failure-prone state that were made, how many were correct. In general, the higher the precision the better

the predictor. Precision is expressed as:

$$Precision = \frac{TP}{TP + FP} \in [0, 1]$$

Recall is the ratio of correctly predicted failures to the number of true failures. In other words, it reports, out of the actual failures that occurred, how many the predictor classified as failure-prone. In conjunction with a higher precision, higher recall is indicative of a better predictor. Recall is expressed as:

$$Recall = \frac{TP}{TP + FN} \in [0, 1]$$

F-Measure, as defined by [18], is the harmonic mean of precision and recall and represents a trade-off between the two. A higher F-Measure reflects a higher quality predictor. F-Measure is expressed as:

$$F-Measure = \frac{2 \cdot Precision \cdot Recall}{Precision + Recall} \in [0, 1]$$

False Positive Rate and Specificity:.

Precision and recall do not account for true negatives (correctly predicted non-failure-prone situations) which can bias an assessment of a predictor. The following measures of performance take true negatives into account to help evaluators more accurately assess and compare predictors.

False Positive Rate (FPR) is the number of incorrectly predicted failures over the total number of predicted non-failure-prone states. A smaller FPR reflects a higher quality predictor. The False Positive Rate is expressed as:

$$FPR = \frac{FP}{FP + TN} \in [0, 1]$$

Specificity the number of times a predictor correctly classified a state as non-failure-prone over all non-failure-prone predictions made. In general, specificity alone is not very useful since failure is rare. Specificity is expressed as:

$$Specificity = \frac{TN}{FP + TN} = 1 - FalsePositiveRate$$

Negative Predictive Value (NPV) and Accuracy:.

In some cases, we wish to show that a prediction approach can correctly classify non-failure-prone situations. The following measures of performance usually can not stand alone due to the nature of failures being rare events. In other words, a highly “accurate” predictor could classify a state 100% of the time as non-failure-prone and still fail to predict every single true failure.

Negative Predictive Value (NPV) is the number of times a predictor correctly classifies a state as non-failure-prone to the total number all non-failure-prone states during which a prediction was made. Higher quality predictors have high NPVs. The NPV is expressed as:

$$NPV = \frac{TN}{TN + FN}$$

Accuracy is the ratio of all correct predictions to the number of predictions made. Accuracy is expressed as:

$$Accuracy = \frac{TP + TN}{TP + FP + FN + TN}$$

Precision/Recall Curve:.

Much like with other predictors, many OFP approaches implement variable thresholds to sacrifice precision for recall or vice versa. That trade-off is typically visualized using a precision/recall curve as seen in Figure 4.




Figure 4. Sample precision/recall curves [19]. Curve *A* represents a poorly performing predictor, curve *B* represents an average predictor, and curve *C* represents an exceptional predictor.

Another popular visualization is the receiver operating characteristic (ROC) curve. By plotting true positive rate over false positive rate one is able to see the predictors ability to accurately classify a failure. A sample ROC curve is shown in Figure 5.




Figure 5. ROC plots of perfect, average, and random predictors [19].

The ROC curve relationship can be further illustrated by calculating the area

under the curve (AUC). Predictors are commonly compared using the AUC which is calculated as follows:

$$AUC = \int_0^1 tpr(fpr) d fpr \in [0, 1],$$

where tpr = true positive rate (recall), and fpr = false positive rate. A pure random predictor will result in an AUC of 0.5 and a perfect predictor a value of 1. The AUC can be thought of as the probability that a predictor will be able to accurately distinguish between a failure-prone state and a non-failure-prone state, over the entire operating range of the predictor.

2.2 Approaches to Online Failure Prediction (OFP)

OFP Taxonomy.

Taxonomies offer natural ways of organizing information and due to the significant body of work surrounding OFP, one was published by Salfner et al. in 2010 [19] classifying many of the OFP approaches in the literature into four major categories. As stated in Section 2.1, these four major categories are defined by the four techniques used to detect faults in real-time: auditing, monitoring, reporting, and tracking. In this research, we update the *reporting* category by classifying approaches to OFP published since its inception. Since this research focusses on real-time *data-driven* device failure prediction approaches, our focus is on the *reporting* category of Salfner’s taxonomy. The *reporting* category organizes failure prediction techniques that attempt to classify a state as failure prone based on reported errors. Salfner et al. further organize the reporting category into five sub-categories: rule-based systems; co-occurrence; pattern recognition; statistical tests; classifiers and they are defined as follows:

Rule-Based Systems attempt to classify a system as being failure-prone or not based a set of conditions met by reported errors. Since modern systems are far too complex to build a set of conditions manually, these approaches seek to find automated ways of identifying these conditions in training data. *Co-occurrence* predictors generate failure predictions based on the reported errors that occur either spatially or temporally close together. *Pattern Recognition* predictors attempt to classify patterns of reported errors as failure prone. This research focusses on pattern recognition OFP approaches, which can be visualized in Figure 6. *Statistical Tests* attempt to classify a system as failure-prone based on statistical analysis of historical data. For example, if a system is generating a much larger volume of error reports than it typically does, it may be a sign of pending failure. *Classifiers* assign labels to given sets of error reports in training data and then make failure predictions based on observed labels in real-time data.




Figure 6. How pattern recognition is accomplished in reported errors [19].

Data-Driven Online Failure Prediction.

The survey published by Salfner et al. covered approaches in every sub-category of the *reporting* category. Since the publication of the survey, we found approaches in two of the subcategories, *pattern recognition* and *classifiers*. We therefore only cover the approaches in those sub-categories of the reporting category here. We found some of the approaches published since Salfner’s survey to be difficult to classify because they employ aspects of the other sub-categories in the *reporting* category. More specifically, many of the modern techniques seem to be a blend between the two

sub-categories *pattern recognition* and *classifiers*. We believe these categories have been blended because these approaches seem to follow general human intuition when looking for software failures. In other words, we have found that cyber operators tend to look for patterns in reported errors and then classify a situation based on those patterns. We therefore categorize these approaches as *hybrid* approaches.

Pattern Recognition:.

In 2006, Salfner et al. proposed an approach to predicting failures by learning patterns of similar events using a semi-Markov chain model [21]. The model learned patterns of error reports that led to failure by mapping the reported errors to the states in the Markov chain and predicted the probability of the transition to a failure-prone state. They tested the model using performance failures of a telecommunication system and reported a precision of 0.8, recall of 0.923, and an F-measure of 0.8571, which drastically outperformed the models to which it was compared.

Given the results, the semi-Markov Chain model is compelling however, it depends on the sequence of reported errors to remain constant in order to be effective. Today, most software is multi-threaded or distributed so there is no guarantee that the sequence of reported errors will remain constant. Further, the authors reported that this approach did not scale well as the complexity of the reported errors grew.

In 2007, Salfner et al. extended their previous work in [21] using semi-Markov models [20]. They generalized the Hidden Semi-Markov process for a continuous-time model and called it the “Generalized Hidden Semi-Markov Model (GHSMM)” By making this generalization, the model was able to effectively predict the sequence of similar events (or in this case, errors) in the continuous time domain. The authors then tested the model and training algorithm using telecommunication performance failure data and compared it to three other approaches. While this GHSMM model

did not perform as well as their previous work, it did outperform the models to which it was compared and more importantly did not depend on the sequence of reported errors. In other words, this new GHMM model predicted failure for permutations of a known failure-prone sequence making it more suited for a distributed or parallel system.

The GHMM approach has been well received by the community, although appears to be limited in use to a single system. Unfortunately, this approach as well as its predecessor, does not scale well and does not adapt to changes to the underlying system without retraining.

Classifiers:.

In 2002, Domeniconi et al. [6] published a technique based on support vector machines (SVM) to classify the present state as either failure prone or not based on a window of error reports as an input vector. As Salfner points out in [19], this SVM approach would not be useful without some sort of transformation of the input vector since the exact same sequence of error messages, rotated by one message, would not be classified as similar. To solve this permutation challenge, the authors in [6] used singular value decomposition to isolate the sequence of error reports that led to a failure.

This SVM approach used training data from a production computer environment with 750 hosts over a period of 30 days. The types of failures the system was trying to detect was the inability to route to a web-page and an arbitrary node being down. Many approaches involving SVMs have been explored since and seem to be popular in the community [8, 9, 16, 6, 13].

Hybrid Approaches:

Since 2012, *Fujitsu Labs* has published several papers on an approach for predicting failure in a cloud-computing environment [22, 25, 24]. Watanabe et al. report on findings after applying a Bayesian learning approach to detect patterns in similar log messages [24, 25]. Their approach abstracts the log messages by breaking them down into single words and categorizing them based on the number of identical words between multiple messages. This hybrid approach removes the details from the messages, like node identifier, and IP address while retaining meaning of the log message.

Watanabe et al.’s hybrid approach attempts to solve the problem of underlying system changes by learning new patterns of messages in real-time. As new messages come in, the model actively updates the probability of failure by Bayesian inference based on the number of messages of a certain type that have occurred within a certain time window. The authors claim that their approach solves three problems: 1) The model is not dependent upon a certain lexicon used to report errors to handle different messages from different vendors; 2) The model does not take into account the order of messages necessarily so in a cloud environment where messages may arrive in different orders, the model is still effective; and 3) The model actively retrain itself so manual re-training does not need to occur after system updates. The model was then tested in a cloud environment over a ninety day period. The authors reported a precision of 0.8 and a recall of 0.9, resulting in an F-measure of 0.847.

In 2012, Fronza et al. [8] introduced a pattern-recognition/classifier hybrid approach that used an SVM to detect patterns in log messages that would lead to failure. The authors used random indexing to solve the problem previously discussed of SVMs failing to classify two sequences as similar if they are offset by one error report. The authors report that their predictor was able to almost perfectly de-

tect non-failure conditions but was poor at identifying failures. The authors then weighted the SVMs to account for this discrepancy by assigning a larger penalty for false negatives than false positives and had better results.

Industry Approaches to Online Failure Prediction.

We thought it necessary to add a brief discussion on how industry seems to have addressed the failure prediction dilemma in recent years. Because hardware has become so easy to acquire, industry seems to have avoided the problem by implementing massive redundancy in their systems. The work in [24, 13] have attributed this problem avoidance to the fact that until recently, implementing and maintaining a failure predictor was difficult. As we decrease the length of the software development life cycle, software updates are being published with increasing frequency leading to rapid changes in underlying systems. These changes can often render a predictor useless without re-training, which is often a manual and resource intensive process.

Adaptive Failure Prediction (AFP) Framework.

The Adaptive Failure Prediction (AFP) Framework by Irrera et al. in [13] seen in Figure 7 presents a new approach to maintaining the efficacy of failure predictors given underlying system changes. The authors conducted a case study implementing the framework using virtualization and fault injection on a web server.

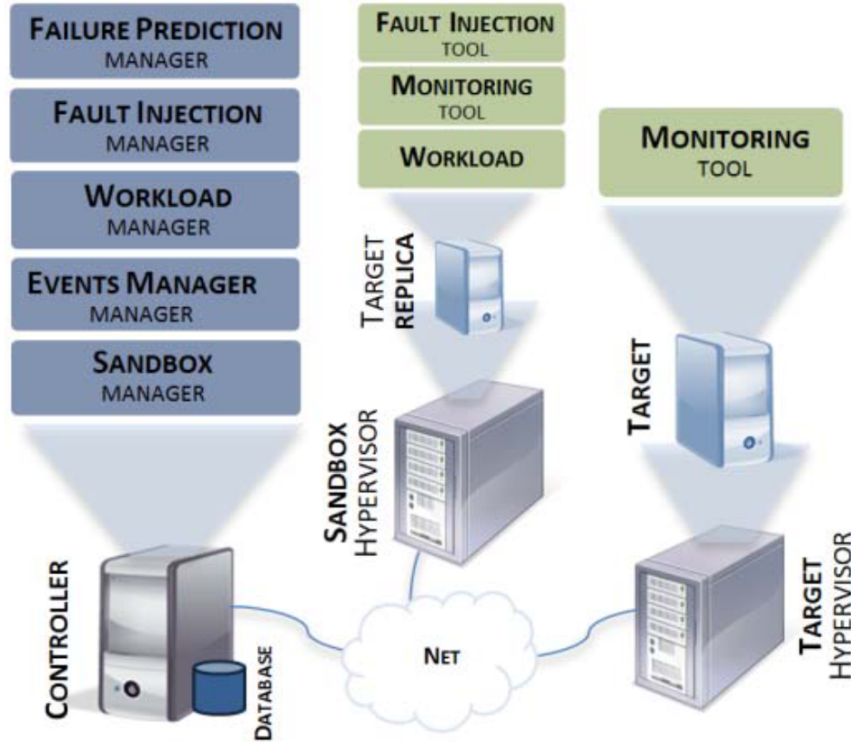


Figure 7. How the AFP framework is implemented [13].

The concept reported used past work by Irrera et al. in [10, 12] to generate failure data by injecting software faults using a tool based on G-SWFIT [7] in a virtual environment for comparing and automatically re-training predictors. In general, the use of simulated data is not well received by the community, however the authors in [11, 12] report evidence supporting the claim that simulated failure data is representative of real failure data. Further, the authors suggest that since systems are so frequently updated and failures are in general rare events, real failure data is often not available. Moreover, the literature shows that even if there is a certain type of failure in training data and a predictor can detect and predict that type of error accurately, it will still miss failures not present in the training data. By injecting the types of

faults that one can expect, each failure type will be represented in the training data.

The authors then conducted a case-study using a web server and an SVM predictor, and report their findings demonstrate their framework is able to adapt to changes to an underlying system which would normally render a predictor unusable.

2.3 Summary

This chapter covered the definitions, measures of performance, and approaches that are relevant to this research as organized under the subsection of *reporting* within the OFP field of study. There has been a tremendous amount of research surrounding the topic of OFP and many prediction approaches have been presented. Unfortunately, these approaches do not appear on modern operational systems and failures are still relatively prevalent. Recent approaches as covered here have sought to make predictors more adaptive to the changes in underlying systems in an effort to make implementing existing failure predictors easier. In this work, we plan to extend the adaptive failure prediction framework and further generalize the approach.

III. Methodology

This chapter outlines the implementation and extensions to the Adaptive Failure Prediction (AFP) Framework as seen in [13] as well as an experiment to validate those extensions and further generalize the framework. The AFP was originally tested on a single system running an operating system that has been deprecated. Consequently, the results from the case study conducted using the AFP are limited in utility and require generalization to be useful to the general community.

The purpose of the AFP framework is to automate the generation of realistic labelled failure data for the purposes of automatically training a failure prediction algorithm. The framework breaks down into modules so that it can be more easily adapted for different applications. This chapter presents three topics. The first describes the process that the extended framework executes in order to generate the labelled training data and automatically train a failure prediction algorithm. The second outlines each module of the extended AFP framework and their associated extensions in detail. The final section outlines extensions to the AFP not covered in the other two sections.

3.1 Failure Data Generation

This work generalizes the Adaptive Failure Prediction framework outlined in [13] by conducting another case study with a Microsoft Windows Server acting as an active directory service. The case study is then taken a step further by showing that the approach holds for other predictors (maybe). Specifically, findings are reported after implementing this framework with the predictor used by Irrera, et al. [13] in their case study and the predictor used by Watanabe, et al. [24].

This section outlines the step-by-step procedure by which the extended AFP is

evaluated to show how effective it is when used on Windows Server deployments. This is done by dividing the steps taken in an experiment into the three major phases as defined in [13]: preparation phase, execution phase, and training phase.

Preparation Phase.

In this phase the AFP is prepared to run for the first time as described in [13]. The Cross Industry Standard Process for Data Mining (CRISP-DM) [4] should be applied to this situation when evaluating how to best apply the AFP for a particular target. For the purposes of this research, our focus is on the Microsoft Windows Directory Services and predicting failure in those services. To demonstrate the efficacy of the AFP, a predictor must be evaluated before and after a significant software update. As a result, the most critical preparation made in evaluating this framework is to hold back all software updates on the target system prior to the first run of the execution phase.

This phase is ultimately the implementation of the framework. Each module of the implementation for this work is detailed in Section 3.2 and is therefore not discussed further here.

Execution Phase.

A general outline of this process can be seen in Figure 8. This phase is divided into three major states: data collection and failure prediction, event checking, and training/update as described in this section.

Data Collection and Failure Prediction.

In this phase, the system has a working predictor providing input to some sort of decision system. It should be noted here that this decision system does not have to be

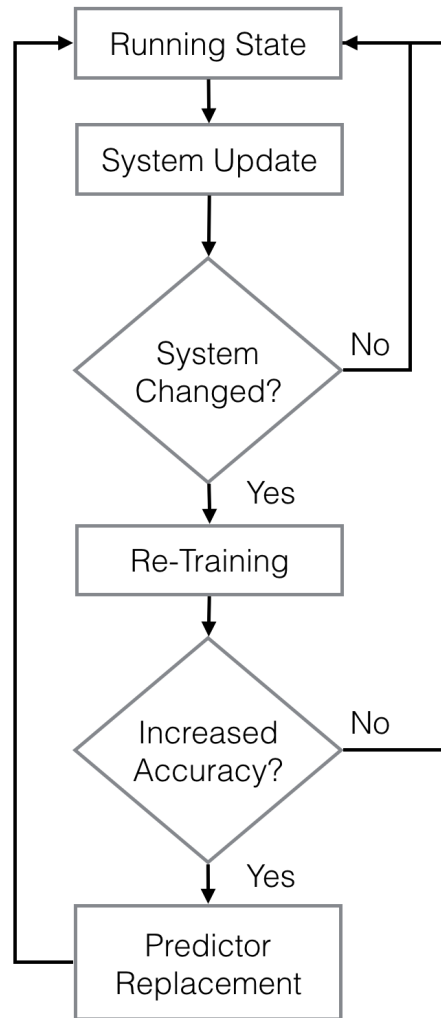


Figure 8. The flow of the major steps involved in the AFP framework execution phase [13].

automated. The system in this phase is making failure predictions about the current state based on the last run of the training phase. For the Air Force application, an operator will likely be the decision maker and the output of the AFP will be a message to that operator.

Event Checking.

Concurrent with the data collection and failure prediction sub-phase, the AFP will be continuously monitoring events that may alter the underlying system. For this experiment, these events are software updates. The output of each episode of this phase is a binary decision to either begin the training phase, or not. In this experiment, the training phase is manually triggered upon completion of a major software update.

Failure Predictor (Re-)Training and Update.

The purpose of this sub-phase is to initiate the training phase and compare its results (a new predictor) with the currently employed predictor. Should the new predictor perform better, the old predictor will be replaced by the new.

Training Phase.

The training phase is broken down into five major steps: target replication, data generation & collection, dataset building, predictor training, and analysis. The general flow can be seen in Figure 9. Each phase is outlined in the following sub-sections.

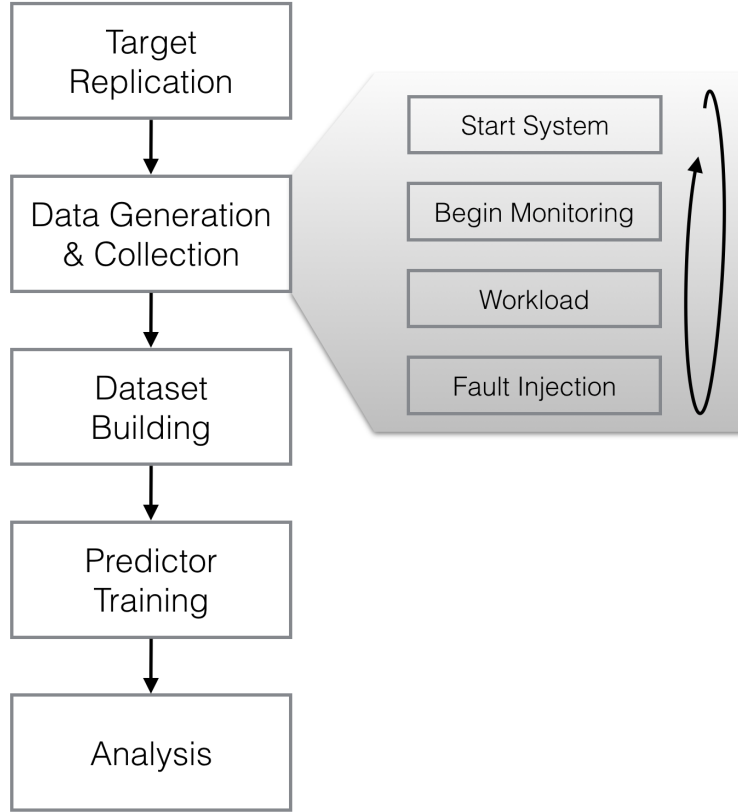


Figure 9. The flow of the major steps involved in the AFP framework training phase [13].

Target Replication.

During this phase a virtual clone of the target is made. After the clone is made, the fault injection and monitoring software must be installed. In this experiment, the monitoring tool is the same as the production system but care must be taken to ensure the hostname is changed so the log messages generated during this phase are not confused with messages from the production system.

Data Generation & Collection.

The purpose of this phase is to generate the data to train a new prediction algorithm. As a result, this sub-phase must be executed several times to generate

statistically acceptable datasets. In this phase, the controller will trigger the cloned target to startup. Once startup is complete and the system enters an idle state, the monitoring tool begins collecting data from the target. After monitoring has begun, the workload is started. Once the workload has entered a steady state, the fault load is started. Finally, when failure occurs, monitoring stops, the workload stops, and the system is rebooted for the next run. To generate golden data, the first run omits the fault injection step.

The most critical part of this process is labelling the data when failure occurs. For the purposes of this experiment, D-PLG reports when authentication fails and transmits a syslog message to the controller. For this research, failure has been defined by two criteria, the first is when authentication with known good credentials fails. The second is classified by when authentication with known good credentials takes longer than thirty seconds (and it can then be assumed that the service has crashed).

Dataset Building.

In this phase, the raw syslog messages are formatted and encoded to train the predictor. The messages generated by the load generator are removed and in each case the neighboring message with the closest timestamp is labelled as when the failure occurred.

TODO: Describe process for encoding messages for training.

Predictor Training.

The purpose of this phase is to use the data generated by the forced failure of the virtual clone to train a machine learning algorithm to classify a system as failure prone or not. This phase is broken into three sub-steps. First, the prediction algorithm is trained using a subset of the generated data. The held-out data from that set is then

used to test the prediction algorithm. Finally, the output of the algorithm is collected for analysis and measure of performance calculation in the next step.

Analysis.

During this phase, the measures of performance in Chapter II are calculated so that the new predictor can be compared against the old. In this experiment, the focus is on the precision and recall of the algorithm. Since failure is such a rare event, accuracy is not a very meaningful measure of performance.

3.2 Implementation of the AFP

AFP Framework Implementation.

This experiment replicates the experiment in [13] except in place of the web-server a Microsoft (MS) Windows Server running Active Directory (AD) Domain Services. In addition, several extensions to the original experiment are made and presented here. Multiple prediction techniques have been applied using this framework to further generalize and validate the framework. The original AFP architecture is shown in Figure 10 with the parts that are modified in this work highlighted.

AFP Modules.

In [13], the authors outline multiple modules into which they have broken the AFP Framework for organizational purposes. This research does not modify these modules, instead, it takes a more granular approach and presents a modified architecture and details each element of that architecture.

The following sections detail the virtual environment in which this architecture was constructed. For reference, this virtual environment was hosted on two VMWare ESXi 5.5 hypervisors each with two 2.6 GHz AMD Opteron 4180 (6 cores each) CPUs

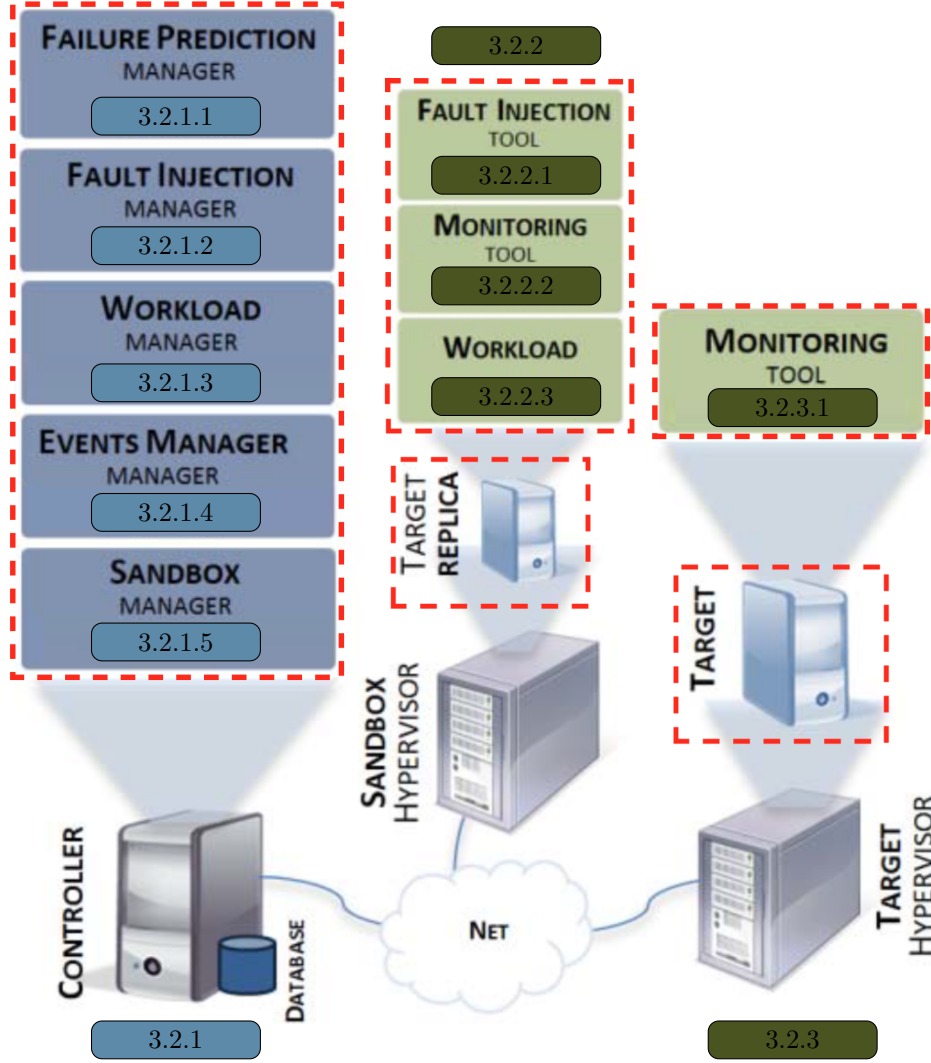


Figure 10. The AFP framework implementation [13] with modified components highlighted.

and 64 GB memory. The individual virtual machines are described in Tables 1, and 2.

3.2.1 Controller Hypervisor.

The controller functions in this experiment are split between two systems on a single hypervisor seen on Table 2. One system is a Microsoft Windows Server responsible for workload management and fault injection management. The additional

Table 1. Hypervisor 1 Configuration (Sandbox/Target)

Qty.	Role	Operating System	CPU / Mem.
1	DC	Win. Server 2008 R2	2 / 2 GB
5	Client	Win. 7	1 / 512 MB

Table 2. Hypervisor 2 Configuration (Controller)

Qty.	Role	Operating System	CPU / Mem.
1	RDP	Win. Server 2008 R2	1 / 4 GB
1	Log	Ubuntu 14.04 LTS	1 / 1 GB

Windows server also hosts remote desktop services to allow the load generator to execute third party authentication with the domain controller. The other system is an Ubuntu 14.04 LTS server that performs the failure prediction management and event management. Each of these functions is detailed in the following sections.

3.2.1.1 Failure Prediction.

The failure prediction module predicts failure using machine learning algorithms trained using the labelled training data generated by the rest of this framework. This module is constantly either training a new predictor because a software update occurred, or predicting failure based on log messages and other features produced by the production system.

The AFP failure prediction function as outlined in [13], is performed by a support vector machine (SVM) predictor using *libsvm*. Additionally, the original experiment made use of a database that stored the features and observations used for the failure prediction training algorithm. This experiment does not modify the failure prediction module drastically as it has already been shown in previous work that the online failure prediction area of study is well explored [19]. This research makes use of a different tool-set to execute the training and predicting phases. Due to its widespread use in the statistical community, the prediction and training algorithms make use of

the R programming language.

In this experiment, an SVM predictor is trained as done in [13], as well as a...
TODO maybe the method in [24]?

3.2.1.2 Fault Injection.

This module is responsible for managing the fault injector installed on the clone of the production target system. The purpose of the fault injector is to force a loaded software application to fail in a realistic way so that the indicators of that failure can be used to train the failure prediction algorithm.

Irrera, et al. [13] use a tool implementing the G-SWFIT technique for this module. G-SWFIT was developed at the University of Coimbra in Coimbra Portugal by Joao Duraes and Henrique Madeira [7]. The method is widely implemented for use in software fault injection both commercially and academically [17, 12, 5, 23]. Recently, studies have questioned the representativeness of the failures generated by G-SWFIT [14]. In each case, the workload generated was critical in creating representative faults. This concern has been addressed in this research and is discussed in Section 3.2.1.3.

An additional concern has been that some faults that have been injected by use of the G-SWFIT technique may not elude modern software testing and as a result never actually occur in production software [17]. The recommended remedy is to conduct source code analysis to determine which pieces of code get executed most frequently and avoid fault injection in those areas since they are most likely to be covered by unit tests. Unfortunately, the target is not an open source project and as a result, some of the faults and resulting failures may never happen in a production environment. Fortunately, the tool that has been developed for this research to do software fault injection automatically scans each library loaded by the target executable for fault

injection points and then is capable of evenly distributing the faults it does inject.

This work introduces an x86-64 implementation of the G-SWFIT technique called W-SWFIT for Windows Software Fault Injection Tool. The source code for W-SWFIT has been published as open source on Github¹ so that others may use it for any of the reasons cited in the original G-SWFIT paper [7]. For this research, the original plan was to use the same tool used in [13] for fault injection. Unfortunately, that tool and all prior G-SWFIT implementations were incapable of injecting faults into x86-64 binary executables. Further, many of the commercial products that were evaluated for this research were incapable of dealing with modern address space layout randomization (ASLR). As a result, W-SWFIT was developed for this research and is capable of injecting faults into all user and kernel mode applications on modern MS Windows operating systems.

The key contributions of W-SWFIT are ASLR adaption and the x86-64 translation we have performed. G-SWFIT works by scanning binary libraries already in memory for patterns (or operators) that match compiled errors in software development. The faults were based on the Orthogonal Defect Classification [2] and can be seen in 3. As pointed out in [19] and [7], failures are ultimately the result of software developer errors. Unfortunately, much of the work done in [7] was on encoding common development errors as IA32 assembly instructions so that working binary executable code could be mutated in memory to introduce these errors in running applications. The target application in this research is strictly an x86-64 (also known as x64 or amd64) application and the patterns identified in [7] are incompatible. Consequently, a fault injection tool capable of mutating x86-64 instructions in the same way was required. W-SWFIT implements many of the operators in [7] in the x86-64 language by translating the operators seen in Table 3 from IA32 to x86-64. A simple example

¹<https://github.com/paullj1/w-swfit/>

of this translation is shown on the entry/exit points of a function in Tables 4, and 5. The rest of the translations can be seen in source code for W-SWFIT. In many cases, the translation was very simple, but in others, the IA32 patterns did not cleanly map to x86-64 byte code. When this happened, great care was taken to ensure the pattern was correctly mapped to x86-64.

Table 3. Table of Faults Injected [7]

Type	Description	ODC Classes
MIFS	Missing "If (cond) { statement(s) }"	Algorithm
MFC	Missing function call	Algorithm
MLAC	Missing "AND EXPR" in expression used as branch	Checking
MLPC	Missing small and localized part of the algorithm	Algorithm
WVAV	Wrong value assigned to a value	Assignment
MVI	Missing variable initialization	Assignment
MVAV	Missing variable assignment using a value	Assignment
WPFV	Wrong variable used in parameter of function call	Interface

Table 4. Funtion Entry/Exit Patterns (IA32) [7]

Module Entry Point		Module Exit Point	
Instruction Sequence	Explanation	Instruction Sequence	Explanation
push ebp	stack frame	move esp,ebp	stack frame
mov ebp, esp	setup	pop ebp	cleanup
sub esp, immedi		ret	

Table 5. Funtion Entry/Exit Patterns (x86-64) [7]

Module Entry Point		Module Exit Point	
Instruction Sequence	Explanation	Instruction Sequence	Explanation
push rbp	stack frame	add rsp, immedi	stack frame
sub rsp, immedi		pop rbp	cleanup
mov rbp, rdx	setup	ret	

3.2.1.3 Workload Managment.

The Workload module creates realistic work for the target system in the sandbox hypervisor to do as a way of generating computational load. Without this module, it

could take a very long time for an injected fault to manifest itself as a failure. Consider a missing *free* statement and the consequent memory leak. A production target server may have a considerable amount of memory and the leak could be very small. To accelerate the possibility of failure occurring, realistic load must be generated against the sandbox clone of the production target.

In the original AFP case study, a Windows XP based web-server was used for a target and therefore the load generation was done by a simple web request generator. As previously mentioned, realistic workload is critical in generating realistic failure and consequently training a useful predictor. As a result, much emphasis is been placed on this module. Since the target is not a web server, it was not possible to use the same load generator as was used in [13]. Initial searches for a load generator suitable for this research yielded a tool developed by Microsoft that initiated remote desktop connections to aid in sizing a terminal services server². By executing a remote desktop session, the authentication and DNS functions of the domain controller would also be loaded. Unfortunately, this tool is no longer maintained and would not execute on the target machine³. Further searches for tools that would sufficiently load the domain controller did not produce any results. Consequently, a tool to produce realistic load for a domain controller was developed for this research and is introduced here.

The Distributed PowerShell Load Generator (D-PLG) is a collection of Microsoft PowerShell scripts designed to generate realistic traffic that will sufficiently load a Microsoft domain controller. Other network traffic generators typically work by replaying traffic captured on a live network. Unfortunately, due to the cryptographic nature of authentication, simply replaying traffic will not load a service since the

²<http://www.microsoft.com/en-us/download/details.aspx?id=2218>

³<https://social.technet.microsoft.com/Forums/windowsserver/en-US/2f8fa5cf-3714-4eb3-a895-c30e2b26862d/debug-assertion-failed-sockcorecpp-line-623>

timestamps and challenge responses will no longer be valid. As a result, any re-played traffic will be dropped and ignored by a live domain controller. D-PLG solves this problem by making native authentication requests by use of built-in PowerShell cmdlets (command-lets). By doing this, realistic authentication requests are sent to the domain controller and are actually processed. The functions performed by the domain controller have been evaluated and D-PLG is designed to sufficiently load each of the services responsible for performing those functions. functions.

In this experiment, the DC is configured as it is in the Air Force Network architecture. After careful analysis, it has been determined that the major roles in the Air Force architecture being performed are authentication and domain name service (DNS). Additionally, by use of native cmdlets, D-PLG is capable of generating four kinds of traffic: web, mail, file sharing, and MS remote desktop protocol (RDP). D-PLG uses the MS Powershell environment to generate the traffic in an effort to make the traffic as real as possible. After building the tool, an experiment was constructed and executed on a scale model of a production environment. The scaled simulation network was built using the recommendations of the Microsoft community for sizing a domain controller [15] and tested by running the tool on five client machines against the domain controller for five rounds of five minutes. The results of this test can be seen in Figures 11, 12, 13.

D-PLG makes use of client machines running a Windows operating system with PowerShell version 4.0 or newer. The controller asks each machine to generate a configurable list of requests at evenly spaced intervals for a configurable duration of time. While this may not be realistic network traffic, it will produce realistic load against a domain controller. Since D-PLG depends on the use of client machines, it is recommended that any load generation be conducted during off-peak hours if spare client sized machines are not available. It should be noted however, that even




Figure 11. How many packets per second were sent or received by the domain controller across all five rounds of the first test. In each test, we captured approximately 1.8 million packets.




Figure 12. How many packets per second were sent or received by one of the clients across all five rounds of the first test.

with poorly resourced client machines (seen in 1), D-PLG was able to generate fifteen thousand authentication sessions over a five minute period; approximately 10 authentication sessions per machine, per second. With modern workstations, the impact on




Figure 13. Domain controller CPU and memory utilization during the first test.

these client machines will likely be negligible and could likely be in use during load generation.

Based on these results, and that a production domain controller should be at approximately 40% CPU utilization during peak utilization [15], we have concluded that D-PLG is capable of sufficiently loading the domain controller over a sustained period of time for the purposes of implementing the AFP framework and is used in this research. Further, we have concluded that D-PLG is capable of scaling to provide load against higher capacity domain controllers by using only a few client machines. There are many more uses for a load generator of this type and we were not able to find a tool that is capable of creating the same type of load so we have published these scripts on Github⁴ for others to use.

3.2.1.4 Events Manager.

This module is responsible for receiving and managing log messages and other events that may be used to train the failure prediction algorithm. Irrera, et al. [13] use the *Logman* tool for event management in their original case study. Since the

⁴<https://github.com/paullj1/AFP-DC/tree/master/D-PLG>

experimental environment was modelled after the Air Force enterprise environment, the *Solar Winds* log forwarding tool is used to perform the functions in this module as it is already present on many of the Air Force domain controllers. The domain controllers on the Sandbox and Target hypervisors forward all events to the Ubuntu virtual machine with the *rsyslog* server daemon configured to receive all messages. These messages are then processed and added to a SQL database for training and prediction.

3.2.1.5 Sandbox Management.

The purpose of the sandbox management module is to supervise the virtual cloning of the production system that is made when a new predictor is to be trained. As Irrera, et al. [13, 10] point out, it is typically inappropriate to inject faults and cause failures in production systems, so a virtual clone must be created for that purpose.

The sandbox is managed manually using Virtual Machine (VM) snapshots. After an initial stable state was configured, snapshots of every component of the architecture were taken so that they could be reset after iterations of the experiment. It is important to note here that because VMWare has documented APIs, in future work, this function could be automated.

3.2.2 Sandbox Hypervisor.

The sandbox is constructed on a single hypervisor implemented as seen on Table 1. The following sections outline each module within this module.

3.2.2.1 Fault Injection.

This module is responsible for causing the target application to fail so that labelled failure data can be generated in a short period of time. As described in Section 3.2.1.2,

W-SWFIT has been developed to serve this purpose and implements the G-SWFIT technique developed by Duraes et al. [7] for fault injection. The execution is controlled by the Windows Server virtual machine on the Controller hypervisor through PowerShell remote execution to reduce the interaction and potential to introduce bias into the training data. The tool allows us to inject a comprehensive list of faults into the AD Services processes and binary libraries which are mostly contained within the 'lsass.exe' process. Since many of the critical functions performed by the AD Services processes are performed in one library called 'ntdsa.dll', it is the focus of fault injection.

3.2.2.2 Monitoring.

The purpose of this module is to capture some evidence or indication of pending failure so that it may be used to train a prediction algorithm. Irrera, et al. [13] use the *Logman* tool in their original study but because the experimental infrastructure used in this research is modelled after production Air Force networks, the *Solar Winds* log forwarding tool is used because it is already present in the Air Force architecture. The tool is a lightweight application that simply forwards windows events to a syslog server.

3.2.2.3 Sandbox Workload.

The sandbox workload module is likely the most critical module in the entire framework. Its purpose is to create realistic work for the target application to do before faults are injected. If the workload is not realistic, then the failures that occur after fault injection will not be representative of real failures and any data or indicators collected cannot be used to train an effective prediction algorithm.

Irrera, et al. [13] used a web traffic generator called TPC-W installed on a single

machine in their original study because their target was a web server. Because the domain controller does not respond to web-requests and a tool had not previously been written for this application, a tool was developed for this research called D-PLG. D-PLG is a tool that generates approximately ten full-stack authentication sessions requests per second in order to sufficiently load the domain controller. D-PLG is a distributed tool and requires the use of client machines as a result. This module is represented by those client machines. In this experiment, the client portion of D-PLG is installed on five client machines managed by the central workload manager as discussed in Section 3.2.1.3.

3.2.3 Target Hypervisor.

The target hypervisor was constructed as a clone of the sandbox hypervisor seen on Table 1. The following section outlines the monitoring tool installed on the domain controller on this hypervisor. It should be noted here that while the client machines were cloned as well for convenience, they were not used in this experiment.

3.2.3.1 Monitoring.

The monitoring module is exactly the same as the sandbox monitoring module and for this experiment, the *Solar Winds* syslog forwarding tool is used. To ensure that the messages that are sent are uniquely identified by the controller, the hostname of the target machine must be different from the hostname of the sandbox target machine.

3.3 Extensions to the AFP

This section outlines a few additional extensions to the Adaptive Failure Prediction Framework. One additional type of failure that may be used to train a predictor

can be generated by under-allocating resources for the domain controller in the sandbox hypervisor. Under some circumstances, this may not be considered a realistic form of failure. However, one reason an organization may want to implement the AFP may be that monetary resources are not available to implement an adequately redundant domain controller and as a result, it may be possible that an adequately sized domain controller is not an option. Consequently, load based failure may be a realistic challenge faced by some organizations and knowing that such a failure may occur might be valuable.

Another type of failure that the extended architecture evaluates goes one layer deeper in analyzing the function performed by the target application. In addition to targeting the main library that performs authentication on a domain controller with fault injection, the extended framework will also target the library responsible for interacting with the user database. In this way, the extended AFP framework is capable of simulating a corrupted database that may occur as a result of a corrupted sector on the disk where the database is stored.

By adding these two additional types of failures to the data used to train a prediction algorithm, the resulting algorithm will be able to predict a wider range of realistic failures.

IV. Experimental Results and Analysis

V. Conclusion and Future Work

Notes on future work:

- More automation using VMWare APIs
- Implement more of the operators from G-SWFIT and better automate injection/training phases
- Automate the event checking process... need a way to better determine when underlying system has changed
- Implement more predictors?

Bibliography

1. A. Avižienis, J. Laprie, B. Randell, and C. Landwehr. Basic concepts and taxonomy of dependable and secure computing. *Dependable and Secure Computing, IEEE Transactions on*, 1(1):11–33, 2004.
2. Norm Bridge and Corinne Miller. Orthogonal Defect Classification Using Defect Data to Improve Software Development. *Software Quality*, 3(1):1–8, 1998.
3. G. Candea, S. Kawamoto, Y. Fujiki, G. Friedman, and A. Fox. Microreboot-a technique for cheap recovery. In *OSDI*, volume 4, pages 31–44, 2004.
4. P. Chapman, J. Clinton, R. Kerber, T. Khabaza, T. Reinartz, C. Shearer, and R. Wirth. CRISP-DM 1.0 Step-by-step data mining guide. Technical report, The CRISP-DM consortium, August 2000.
5. D. Cotroneo, A. Lanzaro, R. Natella, and R. Barbosa. Experimental Analysis of Binary-Level Software Fault Injection in Complex Software. *2012 Ninth European Dependable Computing Conference*, pages 162–172, 2012.
6. C. Domeniconi, C. Perng, R. Vilalta, and S. Ma. A classification approach for prediction of target events in temporal sequences. *Lecture notes in computer science*, pages 125–137, 2002.
7. J. Duraes and H. Madeira. Emulation of software faults: A field data study and a practical approach. *Software Engineering, IEEE Transactions on*, 32(11):849–867, Nov 2006.
8. I. Fronza, A. Sillitti, G. Succi, M. Terho, and J. Vlasenko. Failure prediction based on log files using Random Indexing and Support Vector Machines. *Journal of Systems and Software*, 86(1):2–11, 2013.
9. E. Fulp, G. Fink, and J. Haack. Predicting Computer System Failures Using Support Vector Machines. *Proceedings of the First USENIX conference on Analysis of system logs*, pages 5–5, 2008.
10. I. Irrera, J. Duraes, H. Madeira, and M. Vieira. Assessing the impact of virtualization on the generation of failure prediction data. *Proceedings - 6th Latin-American Symposium on Dependable Computing, LADC 2013*, pages 92–97, 2013.
11. I. Irrera, J. Duraes, M. Vieira, and H. Madeira. Towards Identifying the Best Variables for Failure Prediction Using Injection of Realistic Software Faults. *2010 IEEE 16th Pacific Rim International Symposium on Dependable Computing*, pages 3–10, 2010.

12. I. Irrera and M. Vieira. A Practical Approach for Generating Failure Data for Assessing and Comparing Failure Prediction Algorithms. *2014 IEEE 20th Pacific Rim International Symposium on Dependable Computing*, pages 86–95, 2014.
13. I. Irrera, M. Vieira, and J. Duraes. Adaptive Failure Prediction for Computer Systems: A Framework and a Case Study. *2015 IEEE 16th International Symposium on High Assurance Systems Engineering*, pages 142–149, 2015.
14. N. Kikuchi, T. Yoshimura, R. Sakuma, and K. Kono. Do injected faults cause real failures? A case study of linux. *Proceedings - IEEE 25th International Symposium on Software Reliability Engineering Workshops, ISSREW 2014*, pages 174–179, 2014.
15. S. Makbulolu and G. Geelen. Capacity planning for active directory domain services, 2012.
16. J. Murray, G. Hughes, and K. Kreutz-Delgado. Machine Learning Methods for Predicting Failures in Hard Drives: A Multiple-Instance Application. *J. Mach. Learn. Res.*, 6:783–816, 2005.
17. R. Natella, D. Cotroneo, J. Duraes, and H. Madeira. Representativeness analysis of injected software faults in complex software. *Proceedings of the International Conference on Dependable Systems and Networks*, pages 437–446, 2010.
18. C.J. Rijsbergen. v.(1979). *Information retrieval*, 2, 1979.
19. F. Salfner, M. Lenk, and M. Malek. A survey of online failure prediction methods. *ACM Comput. Surv.*, 42(3):10:1–10:42, March 2010.
20. F. Salfner and M. Malek. Using Hidden Semi-Markov Models for Effective On-line Failure Prediction. *2007 26th IEEE International Symposium on Reliable Distributed Systems (SRDS 2007)*, pages 161–174, 2007.
21. F. Salfner, M. Schieschke, and M. Malek. Predicting failures of computer systems: a case study for a telecommunication system. In *Parallel and Distributed Processing Symposium, 2006. IPDPS 2006. 20th International*, pages 8 pp.–, April 2006.
22. M. Sonoda, Y. Watanabe, and Y. Matsumoto. Prediction of failure occurrence time based on system log message pattern learning. *Network Operations and Management Symposium (NOMS), 2012 IEEE*, (4):578–581, 2012.
23. K. Umadevi and S. Rajakumari. A Review on Software Fault Injection Methods and Tools. pages 1582–1587, 2015.
24. Y. Watanabe. Online Failure Prediction in Cloud Datacenters. 50(1):66–71, 2014.

25. Y. Watanabe, H. Otsuka, M. Sonoda, S. Kikuchi, and Y. Matsumoto. Online failure prediction in cloud datacenters by real-time message pattern learning. *CloudCom 2012 - Proceedings: 2012 4th IEEE International Conference on Cloud Computing Technology and Science*, pages 504–511, 2012.

REPORT DOCUMENTATION PAGE					Form Approved OMB No. 0704-0188	
<p>The public reporting burden for this collection of information is estimated to average 1 hour per response, including the time for reviewing instructions, searching existing data sources, gathering and maintaining the data needed, and completing and reviewing the collection of information. Send comments regarding this burden estimate or any other aspect of this collection of information, including suggestions for reducing this burden to Department of Defense, Washington Headquarters Services, Directorate for Information Operations and Reports (0704-0188), 1215 Jefferson Davis Highway, Suite 1204, Arlington, VA 22202-4302. Respondents should be aware that notwithstanding any other provision of law, no person shall be subject to any penalty for failing to comply with a collection of information if it does not display a currently valid OMB control number. PLEASE DO NOT RETURN YOUR FORM TO THE ABOVE ADDRESS.</p>						
1. REPORT DATE (DD-MM-YYYY)		2. REPORT TYPE		3. DATES COVERED (From — To)		
10-09-2016		Master's Thesis		Sept 2015 — Sep 2016		
4. TITLE AND SUBTITLE DATA DRIVEN DEVICE FAILURE PREDICTION				5a. CONTRACT NUMBER		
				5b. GRANT NUMBER		
				5c. PROGRAM ELEMENT NUMBER		
6. AUTHOR(S) Paul L. Jordan				5d. PROJECT NUMBER		
				5e. TASK NUMBER		
				5f. WORK UNIT NUMBER		
7. PERFORMING ORGANIZATION NAME(S) AND ADDRESS(ES) Air Force Institute of Technology Graduate School of Engineering and Management (AFIT/EN) 2950 Hobson Way WPAFB OH 45433-7765				8. PERFORMING ORGANIZATION REPORT NUMBER AFIT/GCS/ENG/17-M01		
9. SPONSORING / MONITORING AGENCY NAME(S) AND ADDRESS(ES) National Information Assurance Education and Training Program 9800 Savage Road Fort Meade, Maryland 20755-6744 410-854-6206 Email: gmellis@nsa.gov; aeshaff@nsa.gov				10. SPONSOR/MONITOR'S ACRONYM(S) NIETP		
				11. SPONSOR/MONITOR'S REPORT NUMBER(S)		
12. DISTRIBUTION / AVAILABILITY STATEMENT DISTRIBUTION STATEMENT A: APPROVED FOR PUBLIC RELEASE; DISTRIBUTION UNLIMITED.						
13. SUPPLEMENTARY NOTES						
14. ABSTRACT TODO						
15. SUBJECT TERMS Thesis, Failure Prediction, Machine Learning						
16. SECURITY CLASSIFICATION OF:			17. LIMITATION OF ABSTRACT	18. NUMBER OF PAGES	19a. NAME OF RESPONSIBLE PERSON	
a. REPORT	b. ABSTRACT	c. THIS PAGE			Dr. G. Peterson, AFIT/ENG	
U	U	U	U	???	19b. TELEPHONE NUMBER (include area code) (937) 255-????, x????; gilbert.peterson@afit.edu	